

The core software and simulation activities for data analysis at the Pierre Auger Observatory

Eva Santos^{a,*} for the Pierre Auger Collaboration^b

^aFZU - Institute of Physics of the Czech Academy of Sciences, Na Slovance 1999/2, Prague, Czech Republic

^bObservatorio Pierre Auger, Av. San Martín Norte 304, 5613 Malargüe, Argentina Full author list: https://www.auger.org/archive/authors_2024_11.html

E-mail: spokespersons@auger.org

The Pierre Auger Observatory, located near the town Malargüe in the province of Mendoza, Argentina, is the largest cosmic-ray detector in existence, covering an area of 3000 km². The upgraded Observatory, in Phase II of operations, consists of a surface array of 1660 stations combining water Cherenkov, scintillator, and radio detectors. A subset of stations also includes underground muon detectors. Additionally, fluorescence detectors located at four sites overlook the array. The science goals for the enhanced Observatory include the measurement of the properties of ultra-high-energy cosmic rays with large statistics and high sensitivity to the primary composition. The Observatory is also sensitive to photons and neutrinos at the highest energies, allowing it to participate in multi-messenger studies. The Auger Offline Framework provides the tools to perform detailed simulations, using the Geant 4 toolkit, of all components of the Observatory and the analysis of both data and simulated events. It proved to have the flexibility needed to evolve during the lifetime of the Observatory, to accommodate new sub-detectors and, recently, changes to the station readout electronics. A new challenge is interfacing the framework with Machine Learning tools for both the development and execution of neural-network-based algorithms. Independent of the framework, CORSIKA 7 is used to simulate particles, fluorescence light, and radio signals produced by air showers. The production of simulations is coordinated centrally to provide standard libraries for analyses and to optimize the use of computing resources. We will describe the evolution and status of the Offline Framework and the tools used to coordinate the simulation efforts. We will also discuss the challenges of the massive simulation efforts and the resources consumed to provide the simulation libraries required by the Collaboration.

7th International Symposium on Ultra High Energy Cosmic Rays (UHECR2024) 17-21 November 2024 Malargüe, Mendoza, Argentina

^{*}Speaker

1. Introduction

The Auger Offline Framework [1] is one of the keystones of the core software used by the Pierre Auger Collaboration. Its origin dates back to 2002, and, already back then, the Auger Off<u>line</u> was aimed at building software sturdy enough to endure more than 20 years, comprising the whole life of the Pierre Auger Observatory, in which further upgrades would be a reality. The software framework consists of an object-oriented C++-based code designed to be flexible and robust enough to support the simulation and reconstruction of all types of events detected by the Observatory [2]. It can handle several types of data formats and databases, encompassing event simulation and reconstruction, as well as monitoring information from all the instruments. For the event simulation, the reading of the output files from several extensive air shower Monte Carlo codes, namely CORSIKA [3] and AIRES [4], with the options CoREAS [5] and ZHAiRES [6] for the radio component of air showers, and also CONEX [7], for fast simulations of the longitudinal development of air showers, is supported. The advent of AugerPrime [8], the major upgrade of the Pierre Auger Observatory, was the biggest challenge overcome by the Auger Offline Framework, as the implementation of some of its components was found to break the original framework philosophy [9]. Most of the production of the extensive simulation libraries used by a large number of analyses in the Collaboration is done using the grid through the Virtual Organization Auger [10].

2. Core software

The core software of the Pierre Auger Collaboration consists of a set of programs dedicated to the study of high- and ultra-high-energy cosmic rays. At the head is the Auger Offline Framework [1], a modular and universal software used for the reconstruction and simulation of the events detected by the Pierre Auger Observatory, which we detail below. The remaining core software employed by the Pierre Auger Collaboration consists of the Monte Carlo programs CORSIKA [3], AIRES [4], and CONEX [7], which are used to simulate the development of extensive air showers and whose output files are given as input to the Auger Offline Framework for the event simulation. Other programs used are CRPropa [11] and SimProp [12], aimed to simulate the propagation of cosmic rays in an extragalactic environment.

2.1 Auger Offline Framework

The Auger Offline is a universal framework for the reconstruction and simulation of events detected at the Pierre Auger Observatory. It is an object-oriented C++ code with a modular structure aimed at being flexible and robust software, allowing it to be easily extensible to accommodate upgrades to the Pierre Auger Observatory instrumentation throughout its entire lifetime. It was also meant to be a collaborative effort of a large number of physicists developing a variety of applications regarding the reconstruction and simulation of all types of events detected at the Pierre Auger Observatory. For this purpose, the parts of the code directly used by physicists should be clear and relatively straightforward to modify.

2.1.1 The upgraded framework

The initial design of the Pierre Auger Observatory envisaged a hybrid detection technique comprising a 3000 km² surface detector array (SD), composed of 1600 water-Cherenkov detectors

(WCD) arranged in a triangular grid of 1500 m spacing (SD-1500), observed by the *fluorescence* detector (FD), consisting of 24 telescopes placed in four locations on the SD periphery. The first enhancement to the Pierre Auger Observatory targeted an extension to lower energies by deploying 61 additional WCDs in a 23.5 km² denser array of 750 m spacing (SD-750) and HEAT - High Elevation Auger Telescopes, three fluorescence telescopes at the Coihueco site, extending the FD elevation range from 30° to 58°. The Auger Engineering Radio Array (AERA) was later added to the SD array, serving as proof of concept of the feasibility of the future radio detector (RD) of AugerPrime, which incorporates a short aperiodic loaded loop antenna (SALLA) from the RD, and the surface scintillator detector (SSD) atop most WCD. Moreover, the dynamic range of the SD was extended by the installation of a small photomultiplier tube inside each station, and a new electronics board with faster sampling connecting all the detectors was added. Finally, the underground muon detector (UMD), an array of 30 m² area scintillator detectors buried at a depth of 2.3 m at the SD-750 and the newer SD-433 array, in the vicinity of a WCD, aimed at the direct measurement of the muon content of air showers [13, 14] was deployed. While some changes were straightforward to implement, the implementation of AERA required significant additions to the software [15]. Also, some of the AugerPrime components posed challenges to the framework as it meant breaking with some old paradigms of the Auger Offline, called for some unexpected changes, and an update from C++ 98/03 to C++ 11/14 to break strict backward compatibility, eliminate deprecated interfaces, and modernize the development infrastructure [9, 10].

Currently, the Auger Offline can be used to reconstruct events taken during the old (Phase I) and the AugerPrime (Phase II) configurations. The reconstruction and simulation of events comprises those only detected by the SD (SD-1500, SD-750, and SD-433), the FD, and hybrid events, i.e., those detected simultaneously by the SD and the FD, including the HECo extension (HEAT and Coihueco telescopes), and events detected by the AERA, the first radio extension to the Observatory. For Phase II, the novelty is the possibility of the reconstruction of multi-hybrid events, i.e., events detected by the WCD from the SD, including the SSD, the RD, and the UMD.

2.1.2 Structure

The Auger Offline package is organized to have a strict acyclic dependency to avoid problems when building the code. Its structure comprises three main parts, namely, the Modules, the Event, and the Detector, described in detail below. These three components are complemented by a set of foundational classes and utilities for error logging, physics, and mathematical manipulations, as well as packages that support abstract manipulation of geometrical objects. The latter ones do not require any knowledge of the Pierre Auger Observatory or cosmic-ray physics.

Modules The Modules can be assembled and sequenced using instructions in XML format. At the beginning of a run, the framework processes a hierarchy of XML files, starting from the bootstrap file that links all the other necessary XML files for the configuration of the Detector and establishes the module sequence to be executed, see Fig. 1. This XML hierarchy is also used to provide additional configuration data required by the individual modules. The reading and writing operations are handled by special modules, which provide a thin interface to the event input/output (I/O) layer. The names of the files to be processed and written into are part of the data given to the modules.

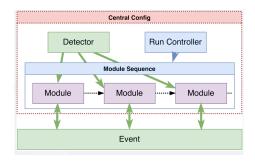


Figure 1: Schematic sketch of the Central Config of the framework defining the Module configuration, sequencing, and the Detector. The Run Controller defines the sequence in which the Modules are executed.

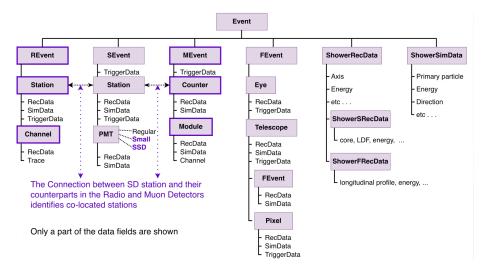


Figure 2: Schematic sketch of the Event data structure. The AugerPrime additions are given in purple. Not all components are shown.

Event The Event serves as a central data structure where modules store and retrieve information. accumulating all reconstruction and, when applicable, also simulation information. It is organized as a collection of classes that mirror the instrument layout of the Pierre Auger Observatory, as well as further subdivisions for accessing Monte Carlo parameters. The Event structure contains all raw, calibrated, reconstructed, and Monte Carlo information, acting as the core framework for communication between modules, as illustrated in Fig. 2. The Event data structure has to be filled with the data expected by the module sequence at the beginning of the Event loop, typically by reading information from one of several available sources. For data events, these formats include the raw event, which is the data format from the DAQ, the internal Off line format, and monitoring formats, where there is information about the raw and calibrated data, as well as the results from the different stages of the event reconstruction. For simulations, the reading of output files from the Monte Carlo codes CORSIKA [3], AIRES [4], and CONEX [7] is supported. In this case, the Event structure also includes information related to detector simulations. Most detailed detector simulations make use of the Geant4 package [16–18] complemented by custom simulations for the tracking of the Cherenkov light, fluorescence photons, and the simulation of the electronics. The Event serialization is implemented using the ROOT toolkit [19]. The internal ROOT-based format

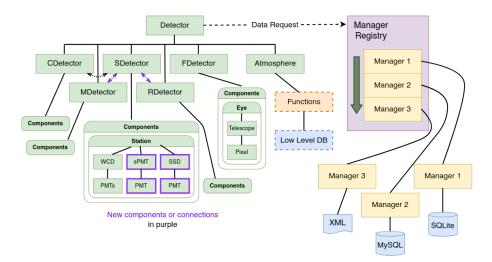


Figure 3: Schematic sketch of the Detector structure, including the access to the atmospheric data. The AugerPrime additions are shown inside the purple boxes.

allows the user to save and restore the complete information in the Event data structure. Similarly, it is also possible to save the result of simulations in the raw data format used for data acquisition or in the *Advanced-Data Summary Tree* (ADST) format [20]. The latter is used for the front-end user analysis based on the I/O component of the ROOT framework, and it is a lightweight tool that can be used in the later stages of data analysis.

Detector The Detector description provides a gateway to data, portraying the detector configuration and performance. Since the atmospheric conditions at the time of the event detection impact the development and properties of extensive air showers, data from the atmospheric monitoring as a function of time are also included as part of the Detector. The Detector components are mapped directly onto data structures. Similarly to Event, the Detector interface follows the hierarchy associated with the instruments of the Pierre Auger Observatory, as shown in Fig. 3. A set of easily usable functions to extract the data is provided. The managers provide access to different data sources; in particular, static detector information is stored in XML format, while time-varying monitoring and calibration data is stored in MySQL and SQLite databases. The Detector structure acts as a facade providing user-visible interfaces to configurable managers, enabling flexible access to various data sources.

2.1.3 Next steps

The increasingly frequent use of *machine learning* (ML) algorithms in a wide range of applications is transforming the traditional way we do scientific analysis, a phenomenon to which the Pierre Auger Collaboration is no exception. See, for instance, [21–24]. In this context, we are implementing a connection between the *Open Neural Network Exchange* [25] (ONNX) ecosystem. This open-source framework provides a standardized representation for ML models and the Auger Offline. Other plans include integrating simulations with the CORSIKA 8 [26] program into the framework.

3. Virtual Organization Auger

The Virtual Organization (VO) Auger, established in 2006, is a group of institutions sharing a common goal and collaborating using distributed computing and data resources provided by the European Grid Infrasture (EGI). The central resources, such as the registration portal and the VOMS (*Virtual Organization Membership Service*) server, are maintained by the CESNET MetaCentrum. All members of the Pierre Auger Collaboration can apply for membership by filling out a registration form, which must be approved by the VO manager. The membership has a validity of approximately one year, after which it can be renewed upon request.

Since 2014, the VO Auger has used the DIRAC (Distributed Infrastructure with Remote Agent Control) interware for job submission, job monitoring, and file catalog management. Our DIRAC server is lodged at the Frances Grilles Infrastructure.

3.1 Grid usage: October 2023 - October 2024

According to the EGI accounting portal, in 2024, the VO Auger ran about 855 thousand single-core jobs on 10 grid sites from 7 countries, totaling 178 million CPU hours, normalized to HEP-Score23 (HS23) [27]. Excluding the VO LHCb from the list of Astrophysics VOs, the VO Auger appears as the second largest EGI user with a relative usage of 17.6%, only behind VIRGO, with 52.4%, as shown in Fig. 4. The VO Auger is consolidating its position as one of the largest users, having increased its HS23

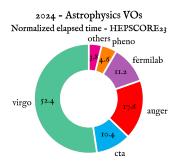
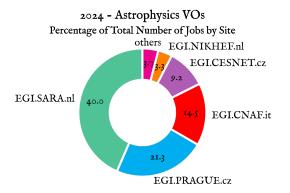


Figure 4: Relative elapsed time of the Astrophysics VOs in 2024, according to the EGI accounting portal. The contribution of the VO LHCb was excluded.

from 77 to 178 million normalized CPU hours in the 2020 - 2024 period. In Fig. 5, the distribution of the relative number of jobs per grid site (left) and the cumulative wall time by site (right) are shown. The bulk of the grid usage in this period comprises the production of CoREAS [5] simu-



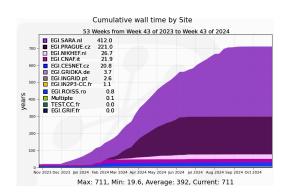


Figure 5: Left: Relative distribution of jobs per grid site. Right: Cumulative wall time per site.

lations. The EGI.SARA.nl site led in both the number of jobs executed and cumulative wall time. Two main factors explain this: only this site accepted long jobs, and possibly the end-of-time of CentOS7, the operating system used to compile CORSIKA and CoREAS considerably worsened

the performance of these jobs. From the experience learned, we advocate for containerizing all software in the CVMFS - *Cern Virtual Machine File System* repository to enhance job performance, particularly in a potentially heterogeneous distributed computing environment.

4. Monte Carlo simulation libraries

The Pierre Auger Collaboration has a task dedicated to the production of the shower and detector official simulation libraries used in a wide variety of physics analyses. Currently, air shower simulations are produced with the CORSIKA [3] program, whose output files are fed into the Auger Offline Framework [1]. So far, the bulk of our detector simulations concern Phase I of the Observatory and comprise SD-only for the SD-1500, SD-750, SD-433, and also hybrid simulations, namely, FD together with SD-1500 and HECo with the SD-750. Special detector simulation libraries comprise hybrid and SD-only time-dependent simulations, which try to reproduce the performance of the SD and FD detectors over time. A detailed description of our shower simulations can be found in [10]. Since then, according to recent findings [28], the cosmic-ray shower library was extended to include Tellurium, Platinum, and Uranium-induced air showers in the energy range 10¹⁷ - 10^{20.2} eV, by using the option CONEX in CORSIKA for the hadronic interaction model EPOS-LHC [29, 30]. Uranium-induced air showers are meant to train ML algorithms on the muon content [23]. Besides the production of the official libraries, the task is also open to simulation requests for specific studies, such as the case of the work reported in [31], or in aiding Collaborators with limited computing resources, as long as their requests are well justified and do not clash with the priorities of the task. The bulk of the simulations are produced using the grid resources provided by the VO Auger. However, for a fraction of CoREAS, whose required CPU time exceeds one week, and other non-standard or test productions, the Prague cluster, which has a maximum wall time of 30 days, is used.

4.1 Accessing the simulations

The preferred method for accessing simulations has been the iRODS (*integrated Rule-Oriented Data System*) provided by the Computing Center IN2P3 in Lyon, France, which will cease its support for the Collaboration at the end of 2025, being replaced by the CNAF Computing Center in Bologna, Italy. Meanwhile, some collaborators have started using the DIRAC Data Management System (DMS), provided they have a valid X509 certificate. DIRAC offers faster file transfers compared to iRODS, as older files at CC IN2P3 are stored on tape, which may significantly increase the download time of each file. However, the issuing of X509 certificates is not possible for all the Collaborators. This problem could be overcome if DIRAC replaces the certification process, such as by issuing tokens for login or making the simulations available in CNAF. Currently, there are about 11 million files with a total size of 1 PB registered in the DIRAC File Catalog.

4.2 Next steps

For 2025, we plan to redo the whole cosmic-ray shower library with the newest CORSIKA 7.7600, which will contain updated hadronic interaction models. The low-energy cut for hadrons will be reduced from 0.05 to 0.02 GeV as it proved to reproduce the late pulses observed by the SSD better [32]. This lower energy cut is enabled by the FLUKA [33, 34] versions 2024.1.0 and

above. The number of proton-induced showers will be increased from 5000 to 10000 as required for signal-to-background studies for the searches of neutral particles, such as photons and neutrinos. We also plan to produce a tau-neutrino-induced air shower library, with and without the radio component of air showers. For long CoREAS simulations, we will exploit the possibility of running MPI jobs on the grid. Finally, we will update our Phase I and Phase II detector simulations using the Auger Offline.

Acknowledgements

This work is co-funded by the Czech Science Foundation under the project GAČR 24-13049S, and by the European Union and by the Czech Ministry of Education, Youth and Sports - Project No. FORTE – CZ.02.01.01/00/22_008/0004632.

References

- [1] S. Argirò et al., Nucl. Instrum. Meth. A 580 (2007) 1485-1496.
- [2] A. Aab et al. [Pierre Auger Collaboration], Nucl. Instrum. Meth. A 798 (2015), 172-213.
- [3] D. Heck et al., Report FZKA 6019 (1998).
- [4] S. J. Sciutto, [astro-ph/9911331].
- [5] T. Huege et al., AIP Conf. Proc., 1535 (2013) 128.
- [6] J. Alvarez-Muniz et al., Astropart. Phys. 35 (2012), 325-341.
- [7] T. Pierog et al., Nucl. Phys. B Proc. Suppl. 151 (2006), 159-162.
- [8] A. Aab et al. [Pierre Auger Collaboration], [arXiv:1604.03637].
- [9] L. Nellen et al. [Pierre Auger Collaboration], PoS ICRC2021 (2021), 250.
- [10] E. Santos et al. [Pierre Auger Collaboration], PoS ICRC2023 (2023), 248.
- [11] R. Alves Batista et al., JCAP 09 (2022), 035.
- [12] R. Aloisio et al., JCAP 11 (2017), 009.
- [13] A. Aab et al. [Pierre Auger Collaboration], JINST 11 (2016) no.02, P02012.
- [14] A. Aab et al. [Pierre Auger Collaboration], JINST 12 (2017) no.03, P03002.
- [15] P. Abreu et al. [Pierre Auger Collaboration], Nucl. Instrum. Meth. A 635 (2011), 92-102.
- [16] V. N. Ivanchenko et al. [Geant4 Collaboration], Nucl. Instrum. Meth. A 502 (2003), 666-668.
- [17] J. Allison et al. [Geant4 Collaboration], IEEE Trans. Nucl. Sci. 53 (2006), 270.
- [18] J. Allison et al. [Geant4 Collaboration], Nucl. Instrum. Meth. A 835 (2016), 186-225.
- [19] R. Brun and F. Rademakers, Nucl. Instrum. Meth. A 389 (1997), 81-86.
- [20] J. G. Gonzalez et al. [Pierre Auger Collaboration], [1208.2154.
- [21] A. Abdul Halim et al. [Pierre Auger Collaboration], Phys. Rev. Lett. 134 (2025) no.2, 021001.
- [22] A. Abdul Halim et al. [Pierre Auger Collaboration], Phys. Rev. D 111 (2025) no.2, 022003.
- [23] S. Hahn et al. [Pierre Auger Collaboration], these proceedings.
- [24] E. Rodriguez et al. [Pierre Auger Collaboration], these proceedings.
- [25] ONNX: https://onnx.ai/; https://github.com/onnx/onnx.
- [26] M. Gottowik et al. [CORSIKA 8 Collaboration], these proceedings.
- [27] R. J. Sobie, [2306.08118].
- [28] B.T. Zhang, these proceedings.
- [29] K. Werner et al., Phys. Rev. C 74 (2006) 044902.
- [30] T. Pierog et al., Phy. Rev. C 92 (2015) 034906.
- [31] F. Ellwanger et al. [Pierre Auger Collaboration], these proceedings.
- [32] T. Schulz et al. [Pierre Auger Collaboration], these proceedings.
- [33] A. Ferrari et al., "FLUKA: A multi-particle transport code", (2005).
- [34] F. Ballarini et al. EPJ Nuclear Sci. Technol. 10 (2024) 16.